

Name - Dilip Puri

ID - 201351014

Collaborator - Hemant Kumar(201352026)

Lab 01

Submission Date - August 13, 2016

Deadline - Due Aug 15, 4:00 PM

Peak Performance

Problem 1 Consider a memory system with a level 1 cache of 32 KB and DRAM of 512 MB with the processor operating at 1 GHz. The latency to L1 cache is 1 cycle and the latency to DRAM is 100 cycles. In each memory cycle, the processor fetches four words (cache line size is 4 words). What is the peak achievable performance of a dot product of two vectors?

```
/* dot product loop */
for (i = 0; i < dim; i++)
    dot_prod += a[i] * b[i];
```

Answer:

```
/* dot product loop */
for (i = 0; i < dim; i++)
    dot_prod += a[i] * b[i];
```

Problem 2 Now consider the problem of multiplying a dense matrix with a vector using a two-loop dot-product formulation. The matrix is of dimension 4K x 4K. (Each row of the matrix takes 16 KB of storage.) What is the peak achievable performance of this technique using a two-loop dot-product based matrix-vector product?

```
/* matrix-vector product loop */
for (i = 0; i < dim; i++)
    for (j = 0; j < dim; j++)
        c[i] += a[i][j] * b[j];
```

Answer:

```
/* dot product loop */

adf
adf
```

Linux commands

```
$top
```




Figure 1: top: provides dynamic real-time view of individual jobs running on the system

```
$gnome-system-monitor
```

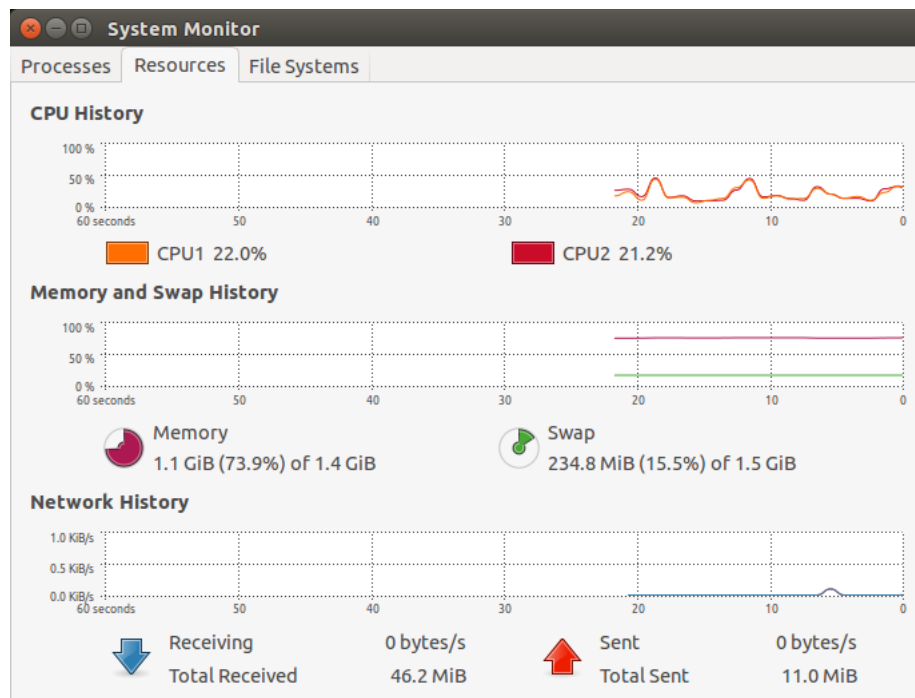


Figure 2: gnome-system-monitor: shows which programs are running and how much processor time, memory, and disk space are being used. This gives an overall system view whereas the “top” instruction represents a detailed perspective

```
$lscpu
```

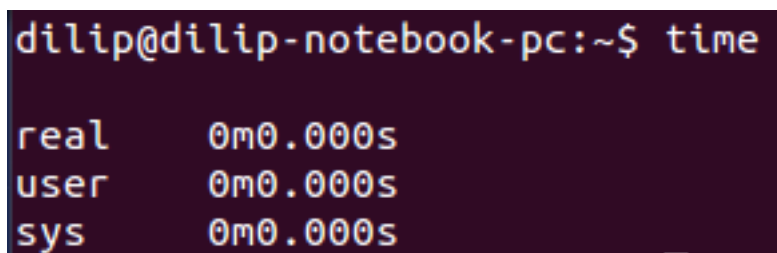
```
We can also get the same information from $cat /proc/cpuinfo
```

```

=====
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                2
On-line CPU(s) list:   0,1
Thread(s) per core:    1
Core(s) per socket:    2
Socket(s):             1
NUMA node(s):          1
Vendor ID:             AuthenticAMD
CPU family:            18
Model:                 1
Stepping:              0
CPU MHz:               800.000
BogoMIPS:              4392.08
Virtualization:        AMD-V
L1d cache:             64K
L1i cache:             64K
L2 cache:              1024K
NUMA node0 CPU(s):    0,1
dilip@dilip-notebook-pc:~$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                2
On-line CPU(s) list:   0,1
Thread(s) per core:    1
Core(s) per socket:    2
Socket(s):             1
NUMA node(s):          1
Vendor ID:             AuthenticAMD
CPU family:            18
Model:                 1
Stepping:              0
CPU MHz:               800.000
BogoMIPS:              4392.08
Virtualization:        AMD-V
L1d cache:             64K
L1i cache:             64K
L2 cache:              1024K
NUMA node0 CPU(s):    0,1

```

```
$time ./a.out
```



```

dilip@dilip-notebook-pc:~$ time

real    0m0.000s
user    0m0.000s
sys     0m0.000s

```

Figure 3: time: Get total program execution time in the shell

Lab Problems

1. Familiarize yourself with the Linux commands and POSIX thread code given in this handout.
2. Solutions for Problems 1-2 on peak performance.
3. Using the basic Linux commands find the cache size, bandwidth number of processors on your system.
4. Write a C-code using POSIX threads to create an unbalanced load using sleep command and hello.c. The sample sleep times are given below:
 - (a) thread-1: 1000 sec, thread-2: 5000 sec, thread-3: 20 sec, thread-4: 1200 sec.
 - (b) Measure the total time taken for the complete execution of code with and without the additional sleep command.Hint: You will require to include unistd.h for successful compilation.



```
hemant@hemant:~/Desktop/sem7/pp/lab$ gcc -pthread -o hello hello.c
hemant@hemant:~/Desktop/sem7/pp/lab$ time ./hello
In main: creating thread 0
In main: creating thread 1
In main: creating thread 2
In main: creating thread 3
Hello World! It's me, thread #2!
In main: creating thread 4
Hello World! It's me, thread #4!
Hello World! It's me, thread #0!
Hello World! It's me, thread #3!
Hello World! It's me, thread #1!

real    0m50.002s
user    0m0.000s
sys     0m0.003s
```

Figure 4: time: Get total program execution time in the shell

5. Write a C-code using POSIX threads about matrix multiplication.
 - (a) Take overall execution time measurement using time command for different application size and thread count for the serial and parallel code.
 - (b) Observe gnome-system-monitor output as your fire up different thread counts.
 - (c) Use the overall execution time measurements to plot and comment upon the speed-up.