

**Name** - Dilip Puri

**ID** - 201351014

**Collaborator** - Hemant Kumar(201352026)

---

### Lab 01

---

**Submission Date** - August 15, 2016

**Deadline** - Due Aug 15, 4:00 PM

---

## Peak Performance

**Problem 1** Consider a memory system with a level 1 cache of 32 KB and DRAM of 512 MB with the processor operating at 1 GHz. The latency to L1 cache is 1 cycle and the latency to DRAM is 100 cycles. In each memory cycle, the processor fetches four words (cache line size is 4 words). What is the peak achievable performance of a dot product of two vectors?

```
/* dot product loop */  
for (i = 0; i < dim; i++)  
    dot_prod += a[i] * b[i];
```

-----  
**Answer:**

Main memory = 512 MB  
Cache = 32 KB  
Processor operating at 1 GHz  
L1 cache latency = 1 cycle  
L2 or main memory latency = 100 cycle  
In 1 memory cycle processor fetch no. of words = 4  
cache line size = 4 words  
Peak Performance = no. of maximum floating pt(arithmetic) operations  
so,

1 cycle takes  $1/(1 \times 10^9)$  seconds.  
One access to memory takes  $100/(1 \times 10^9)$  seconds. = 100ns.  
Performance =  $4 \text{ FLOPS} / (100 / (1 \times 10^9)) = 40 \text{ MFLOPS}$ .

-

**Problem 2** Now consider the problem of multiplying a dense matrix with a vector using a two-loop dot-product formulation. The matrix is of dimension 4K x 4K. (Each row of the matrix takes 16 KB of storage.) What is the peak achievable performance of this technique using a two-loop dot-product based matrix-vector product?

```
/* matrix-vector product loop */
for (i = 0; i < dim; i++)
    for (j = 0; j < dim; j++)
        c[i] += a[i][j] * b[i];
```

-----

Answer:

Main memory = 512 MB

Cache = 32 KB

Dimension of matrix = 4K(16 KB storage)

1 cycle takes  $1/(1 \times 10^9)$  seconds.

One access to memory takes  $100/(1 \times 10^9)$  seconds. = 100ns.

Performance =  $8 \text{ FLOPS}/(100/(1 \times 10^9)) = 80 \text{ MFLOPS}$ .

-

## Linux commands

\$top

```
top - 02:47:51 up 12:02, 2 users, load average: 0.41, 0.43, 0.44
Tasks: 187 total, 2 running, 185 sleeping, 0 stopped, 0 zombie
%Cpu(s): 5.6 us, 2.4 sy, 0.0 ni, 87.7 id, 4.4 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 1510908 total, 1407324 used, 103584 free, 23852 buffers
KiB Swap: 1547260 total, 240476 used, 1306784 free, 303060 cached Mem

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 2081 dillip    20   0 1552804 153936 22560 S  6.6 10.2 11:11.04 compliz
1408 root      20   0 323344 59880 25448 S  6.0  4.0 18:38.43 Xorg
7610 dillip    20   0 631520 16880 12040 S  2.3  1.1  0:00.92 gnome-screensho
5305 dillip    20   0 1465000 459236 41320 S  1.7 30.4 28:09.40 firefox
   8 root      20   0   0   0   0 S  0.3  0.0  0:29.07 rcuos/0
  788 root      20   0 492764 3676 2248 S  0.3  0.2  0:01.61 NetworkManager
1097 mysql     20   0 550096 8144 804 S  0.3  0.5  0:19.51 mysqld
6716 root      20   0   0   0   0 S  0.3  0.0  0:09.86 kworker/0:0
7584 dillip    20   0 651572 19280 12760 S  0.3  1.3  0:00.81 gnome-terminal
   1 root      20   0 33756 2108 836 S  0.0  0.1  0:02.23 init
   2 root      20   0   0   0   0 S  0.0  0.0  0:00.02 kthreadd
   3 root      20   0   0   0   0 S  0.0  0.0  0:00.43 ksoftirqd/0
   5 root      0 -20   0   0   0 S  0.0  0.0  0:00.00 kworker/0:0H
   7 root      20   0   0   0   0 R  0.0  0.0  0:38.26 rcu_sched
   9 root      20   0   0   0   0 S  0.0  0.0  0:28.37 rcuos/1
  10 root      20   0   0   0   0 S  0.0  0.0  0:00.00 rcuos/2
  11 root      20   0   0   0   0 S  0.0  0.0  0:00.00 rcuos/3
  12 root      20   0   0   0   0 S  0.0  0.0  0:00.00 rcu_bh
  13 root      20   0   0   0   0 S  0.0  0.0  0:00.00 rcuob/0
  14 root      20   0   0   0   0 S  0.0  0.0  0:00.00 rcuob/1
  15 root      20   0   0   0   0 S  0.0  0.0  0:00.00 rcuob/2
  16 root      20   0   0   0   0 S  0.0  0.0  0:00.00 rcuob/3
  17 root      rt   0   0   0   0   0 S  0.0  0.0  0:00.56 migration/0
  18 root      rt   0   0   0   0   0 S  0.0  0.0  0:00.17 watchdog/0
  19 root      rt   0   0   0   0   0 S  0.0  0.0  0:00.15 watchdog/1
```

Figure 1: top: provides dynamic real-time view of individual jobs running on the system

\$gnome-system-monitor

\$lscpu

We can also get the same information from \$cat /proc/cpuinfo

```
=====
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:             Little Endian
CPU(s):                 2
On-line CPU(s) list:   0,1
Thread(s) per core:    1
Core(s) per socket:    2
Socket(s):              1
NUMA node(s):          1
```

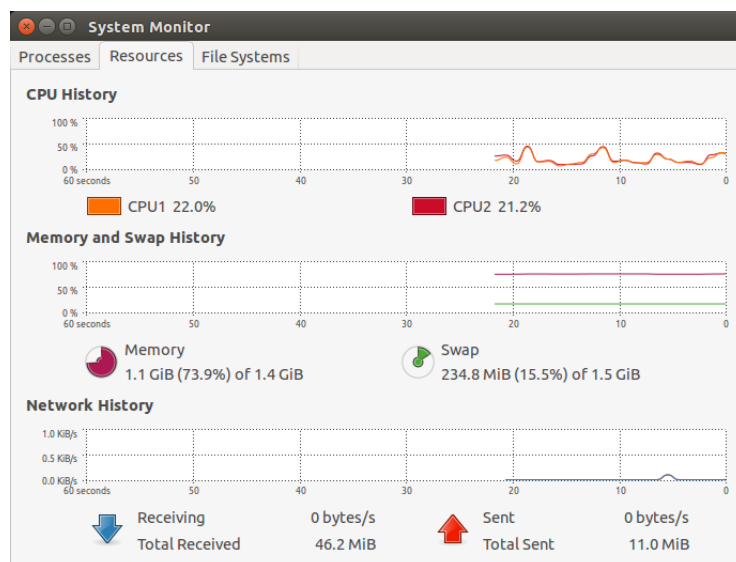


Figure 2: gnome-system-monitor: shows which programs are running and how much processor time, memory, and disk space are being used. This gives an overall system view whereas the “top” instruction represents a detailed perspective

```
Vendor ID: AuthenticAMD
CPU family: 18
Model: 1
Stepping: 0
CPU MHz: 800.000
BogoMIPS: 4392.08
Virtualization: AMD-V
L1d cache: 64K
L1i cache: 64K
L2 cache: 1024K
NUMA node0 CPU(s): 0,1
dilip@dilip-notebook-pc:~$ lscpu
Architecture: x86_64
CPU op-mode(s): 32-bit, 64-bit
Byte Order: Little Endian
CPU(s): 2
On-line CPU(s) list: 0,1
Thread(s) per core: 1
Core(s) per socket: 2
Socket(s): 1
NUMA node(s): 1
Vendor ID: AuthenticAMD
CPU family: 18
Model: 1
Stepping: 0
CPU MHz: 800.000
BogoMIPS: 4392.08
Virtualization: AMD-V
L1d cache: 64K
L1i cache: 64K
L2 cache: 1024K
NUMA node0 CPU(s): 0,1
```

```
$time ./a.out
```

```
dilip@dilip-notebook-pc:~$ time
real    0m0.000s
user    0m0.000s
sys     0m0.000s
```

Figure 3: time: Get total program execution time in the shell

## Lab Problems

1. Familiarize yourself with the Linux commands and POSIX thread code given in this handout.
2. Solutions for Problems 1-2 on peak performance.
3. Using the basic Linux commands find the cache size, bandwidth number of processors on your system.
4. Write a C-code using POSIX threads to create an unbalanced load using sleep command and hello.c. The sample sleep times are given below:
  - (a) thread-1: 1000 sec, thread-2: 5000 sec, thread-3: 20 sec, thread-4: 1200 sec.
  - (b) Measure the total time taken for the complete execution of code with and without the additional sleep command.  
Hint: You will require to include unistd.h for successful compilation.
5. Write a C-code using POSIX threads about matrix multiplication.
  - (a) Take overall execution time measurement using time command for different application size and thread count for the serial and parallel code.
  - (b) Observe gnome-system-monitor output as your fire up different thread counts.
  - (c) Use the overall execution time measurements to plot and comment upon the speed-up.

